

2022 Spring

---

# Databases-1 Lecture-01

Introduction, Relational Algebra

# Information, 2022 Spring

---

- ▶ Databases-1 Lecture:  
Monday 16:00-17:30 online in MS Teams  
Tuesday 12:15-13:45 online in MS Teams
- ▶ Instructor: Tibor Nikovits Dr.  
<http://people.inf.elte.hu/nikovits>
- ▶ Website of the course:  
<http://people.inf.elte.hu/nikovits/DB1>


# Textbook

---

- ▶ A First Course in Database Systems (3rd ed.)  
by Jeff Ullman and Jennifer Widom (UW)

same material and sections as

- ▶ Database Systems: The Complete Book (2nd ed)  
by Garcia-Molina, Jeff Ullman and Jennifer Widom



# Topics of the semester

---

- ▶ Relational Data Model
- ▶ Core and Extended Relational Algebra
- ▶ SQL Query and Modification
- ▶ Constraints, Triggers and Views
- ▶ Persistent Stored Modules, Oracle PL/SQL
- ▶ Datalog, Recursion
- ▶ Entity-Relationship Model
- ▶ Design of Relational Databases

# What is a Data Model?

---

- ▶ 1. Mathematical representation of data
- ▶ 2. Operations on data
- ▶ 3. Constraints

# Relational Data Model

---

- ▶ A relation is a table

Attributes (column headers)

Beer	Manufacturer
Winterbrew	Pete's
Bud Lite	Anheuser-Busch

Tuples (rows)

# Types and schemas

---

- ▶ *Relation schema* = relation name + attributes, in order (+ types of attributes).
  - ▶ Example: Beers(name, manf) or Beers(name: string, manf: string)
- ▶ *Database* = collection of relations.
- ▶ *Database schema* = set of all relation schemas in the database.

# Why relations?

---

- ▶ Very simple model.
- ▶ *Often* matches how we think about data.
- ▶ Abstract model that underlies **SQL**, the most important **database language** today.



# Relational model

---

- ▶ Logical level:
  - ▶ The relations are considered as tables.
  - ▶ The tables have unique names
  - ▶ The columns address the attributes
  - ▶ The rows represent the records
  - ▶ Rows can be interchanged, the order of rows is irrelevant
- ▶ Physical level:
  - ▶ The relations are stored in a file structure

# Examples

---

Example 1

A	B	C
a	b	c
d	a	a
c	b	d

Example 2

B	C	A
b	c	a
a	a	d
d	d	c

In ex. 1 and ex. 2 the columns are interchanged but the same relation

Example 3

A	B	C
c	b	d
d	a	a
a	b	c

Example 4

A	B	C
c	b	d
c	b	d
a	b	c

In ex. 1 and ex. 3 the same tuples are represented in different order, but these are the same relations too.

Ex. 4 is not a relation

# Defining a Database Schema

---

- ▶ A database schema comprises declarations for the relations (“tables”) of the database.
- ▶ Many other types of objects may also appear in the database schema, including views, constraints, triggers, indexes, etc.

# Declaring a Relation

---

▶ Simplest form is:

▶ **CREATE TABLE** <name> (<list of elements>);

```
CREATE TABLE Sells (  
    bar        CHAR(20) ,  
    beer       VARCHAR(20) ,  
    price      REAL  
);
```

# Elements of Table Declarations

---

- ▶ The principal element is a pair consisting of an **attribute** and a **datatype**.
- ▶ The most common datatypes are:
  - ▶ INT or INTEGER (synonyms).
  - ▶ REAL or FLOAT (synonyms).
  - ▶ CHAR( $n$ ) = fixed-length string of  $n$  characters.
  - ▶ VARCHAR( $n$ ) = variable-length string of up to  $n$  characters.
  - ▶ DATE is a type, and the form of a date value is:  
Example: 'yyyy-mm-dd' DATE '2002-09-30'

# Example: Create Table

---

```
CREATE TABLE Sells (  
    bar          CHAR(20) ,  
    beer         VARCHAR(20) ,  
    price        REAL  
);
```

# Other Declarations for Attributes

---

- ▶ Basic declaration for an attribute is a pair consisting of an attribute and a data type.
- ▶ Other declarations we can make for an attribute are:
  1. **NOT NULL** means that the value for this attribute may never be NULL.
  2. **DEFAULT** <value> says that if there is no specific value known for this attribute's component in some tuple, use the stated <value>.

# Example: Default Values

---

```
CREATE TABLE Drinkers (  
    name CHAR(30) NOT NULL,  
    addr CHAR(50) DEFAULT '123 Sesame St.',  
    phone CHAR(16)  
);
```



# Effect of Defaults -- 1

---

- ▶ Suppose we insert the fact that Sally is a drinker, but we know neither her address nor her phone.
- ▶ An INSERT with a **partial list of attributes** makes the insertion possible:

```
INSERT INTO Drinkers (name)
VALUES ('Sally');
```

## Effect of Defaults -- 2

---

- ▶ But what tuple appears in Drinkers?

name	addr	phone
'Sally'	'123 Sesame St'	NULL

- ▶ If we had declared phone NOT NULL, this insertion would have been rejected.

# Remove a relation from schema

---

- ▶ Remove a relation from the database schema by:

- ▶ **DROP TABLE <name>;**

- ▶ Example:

```
DROP TABLE Sells;
```

# Query Languages: Relational Algebra

---

- ▶ What is an “Algebra”?
- ▶ Mathematical system consisting of:
  - ▶ *Operands* --- variables or values from which new values can be constructed.
  - ▶ *Operators* --- symbols denoting procedures that construct new values from given values.

# Core Relational Algebra

---

- ▶ Union, intersection, and difference.
  - ▶ Usual set operations but require that the two operands have the same relation schema.
- ▶ Selection: picking certain rows.
- ▶ Projection: picking certain columns.
- ▶ Products and joins: compositions of relations.
- ▶ Renaming of relations and attributes.

# Union, intersection, difference

---

- ▶ To apply these operators the relations must have the same attributes.
- ▶ Union ( $R1 \cup R2$ ): all tuples from R1 or R2
- ▶ Intersection ( $R1 \cap R2$ ): common tuples from R1 and R2
- ▶ Difference ( $R1 \setminus R2$ ): tuples occurring in R1 but not in R2

# Example

---

Relation Sells1:

Bar	Beer	Price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50

Relation Sells2:

Bar	Beer	Price
Joe's	Bud	2.50
Jack's	Bud	2.75

$\text{Sells1} \cup \text{Sells2}$ :

Bar	Beer	Price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50
Jack's	Bud	2.75

$\text{Sells1} \cap \text{Sells2}$ :

Bar	Beer	Price
Joe's	Bud	2.50

$\text{Sells2} \setminus \text{Sells1}$ :

Bar	Beer	Price
Jack's	Bud	2.75



# Selection

---

- ▶  $R1 := \sigma_C(R2)$ 
  - ▶  $C$  is a condition (as in “if” statements) that refers to attributes of  $R2$ .
  - ▶  $R1$  is all those tuples of  $R2$  that satisfy  $C$ .



# Example

---

Relation Sells:

Bar	Beer	Price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50
Sue's	Miller	3.00

JoeMenu :=  $\sigma_{\text{bar}=\text{"Joe's"}}(\text{Sells})$ :

Bar	Beer	Price
Joe's	Bud	2.50
Joe's	Miller	2.75

# Projection

---

- ▶  $R1 := \pi_L(R2)$ 
  - ▶  $L$  is a list of attributes from the schema of  $R2$ .
  - ▶  $R1$  is constructed by looking at each tuple of  $R2$ , extracting the attributes on list  $L$ , in the order specified, and creating from those components a tuple for  $R1$ .
  - ▶ Eliminate duplicate tuples, if any.

# Example

---

Relation Sells:

Bar	Beer	Price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50
Sue's	Miller	3.00

Prices :=  $\pi_{\text{beer,price}}(\text{Sells})$ :

Beer	Price
Bud	2.50
Miller	2.75
Miller	3.00

# Product

---

- ▶  $R3 := R1 \times R2$ 
  - ▶ Pair each tuple  $t1$  of  $R1$  with each tuple  $t2$  of  $R2$ .
  - ▶ Concatenation  $t1t2$  is a tuple of  $R3$ .
  - ▶ Schema of  $R3$  is the attributes of  $R1$  and  $R2$ , in order.
  - ▶ But take care if an attribute ( $A$ ) *has the* same name in  $R1$  and  $R2$ : use  $R1.A$  and  $R2.A$ .

# Example: $R3 = R1 \times R2$

---

► R1

A	B
1	2
3	4

► R2

B	C
5	6
7	8
9	10

$R3 = R1 \times R2$

A	R1.B	R2.B	C
1	2	5	6
1	2	7	8
1	2	9	10
3	4	5	6
3	4	7	8
3	4	9	10

# Theta-Join

---

- ▶  $R3 := R1 \bowtie_C R2$ 
  - ▶ Take the product  $R1 * R2$ .
  - ▶ Then apply  $\sigma_C$  to the result.
- ▶ As for  $\sigma$ ,  $C$  can be any boolean-valued condition.
  - ▶ Historic versions of this operator allowed only  $A \text{ theta } B$ , where theta was  $=$ ,  $<$ , etc.; hence the name “theta-join.”

# Example

Sells:

Bar	Beer	Price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50
Sue's	Miller	3.00

Bars:

Name	Address
Joe's	Maple st.
Sue's	River rd.

Barinfo= Sells ⋈<sub>Sells.bar = Bars.name</sub> Bars

Bar	Beer	Price	Name	Address
Joe's	Bud	2.50	Joe's	Maple st.
Joe's	Miller	2.75	Joe's	Maple st.
Sue's	Bud	2.50	Sue's	River rd.
Sue's	Miller	3.00	Sue's	River rd.

# Natural Join

---

- ▶ A frequent type of join connects two relations by:
  - ▶ Equating attributes of the same name, and
  - ▶ Projecting out one copy of each pair of equated attributes.
- ▶ Called *natural* join.
- ▶ Denoted  $R3 := R1 \bowtie R2$ .



# Example

---

Sells:

Bar	Beer	Price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50
Sue's	Miller	3.00

Bars:

Bar	Address
Joe's	Maple st.
Sue's	River rd.

Barinfo= Sells ⋈ Bars

Bar	Beer	Price	Address
Joe's	Bud	2.50	Maple st.
Joe's	Miller	2.75	Maple st.
Sue's	Bud	2.50	River rd.
Sue's	Miller	3.00	River rd.

# Renaming

---

- ▶ The RENAME operator gives a new schema to a relation.
- ▶  $R1 := \rho_{R1(A1, \dots, An)}(R2)$  makes R1 be a relation with attributes  $A1, \dots, An$  and the same tuples as R2.
- ▶ **Simplified notation:**  $R1(A1, \dots, An) := R2$ .

# Example

---

Bars:

Name	Address
Joe's	Maple st.
Sue's	River rd.

$R(\text{Bar}, \text{Address}) := \text{Bars}$

Bar	Address
Joe's	Maple st.
Sue's	River rd.

# Building Complex Expressions

---

- ▶ **Algebras** allow us to express **sequences of operations** in a natural way
  - ▶ Example: in arithmetic  $\rightarrow (x + 4) * (y - 3)$ .
- ▶ Relational algebra allows the same.
- ▶ **Three notations**, just as in arithmetic:
  1. **Sequences of assignment** statements.
  2. Expressions with **several operators**.
  3. Expression **trees**.

# Sequences of Assignments

---

- ▶ Create **temporary relation** names.
- ▶ Renaming can be implied by giving relations a list of attributes.
- ▶ Example:  $R3 := R1 \bowtie_C R2$  can be written:  
 $R4 := R1 \times R2$   
 $R3 := \sigma_C(R4)$

# Expressions in a Single Assignment

---

- ▶ Example: the theta-join  $R3 := R1 \bowtie_C R2$  can be written:  $R3 := \sigma_C (R1 \times R2)$
- ▶ **Precedence** of relational operators:
  1. Unary operators --- select, project, rename --- have highest precedence, bind first.
  2. Then come products and joins.
  3. Then intersection.
  4. Finally, union and set difference bind last.
- ▶ But you can always **insert parentheses** to force the order you desire.

# Expression Trees

---

- ▶ Leaves are operands --- either variables standing for relations or particular, constant relations.
- ▶ Interior nodes are operators, applied to their child or children.

# Example


---

- ▶ Using the relations Bars(name, address) and Sells(bar, beer, price), find the names of all the bars that are either on Maple St. or sell Bud for less than \$3.



# As a Tree:

---



# Schema-Defining Rules

---

- ▶ For **union**, **intersection**, and **difference**, the schemas of the two operands must be the same, so use that schema for the result.
- ▶ **Selection**: schema of the result is the same as the schema of the operand.
- ▶ **Projection**: list of attributes tells us the schema.
- ▶ **Product**, **Theta-join**: the schema is the attributes of the two relations.
  - ▶ Use  $R.A$ , etc., to distinguish two attributes named  $A$ .
- ▶ Natural join: use attributes of the two relations.
  - ▶ Common attribute names are merged into one.
- ▶ **Renaming**: the operator tells the schema.

# Relational algebra: Monotony

- ▶ Monotone non-decreasing expression:
  - ▶ applied on more tuples, the result contains more tuples
  - ▶ Formally if  $R_i \subseteq S_i$  for every  $i=1,\dots,n$ , then  $E(R_1,\dots,R_n) \subseteq E(S_1,\dots,S_n)$ .
- ▶ **Difference** is the only operator which is **not** monotone:

A	B		A	B		A	B		A	B
1	0	-	1	0	⊄	1	0	-	1	0
2	1					2	1		2	1