

Introduction to SQL

Select-From-Where Statements

Multirelation Queries

Subqueries

Why SQL?

- SQL is a very-high-level language.
 - Say “what to do” rather than “how to do it.”
 - Avoid a lot of data-manipulation details needed in procedural languages like C++ or Java.
- Database management system figures out “best” way to execute query.
 - Called “query optimization.”

Select-From-Where Statements

SELECT desired attributes

FROM one or more tables

WHERE condition about tuples of
the tables

Our Running Example

- All our SQL queries will be based on the following database schema.

- Underline indicates key attributes.

Beers(name, manf)

Bars(name, addr, license)

Drinkers(name, addr, phone)

Likes2(drinker, beer) → Likes2 !!!

Sells(bar, beer, price)

Frequents(drinker, bar)

Example

- Using **Beers(name, manf)**, what beers are made by Anheuser-Busch?

```
SELECT name
```

```
FROM Beers
```

```
WHERE manf='Anheuser-Busch';
```

Result of Query

name
Bud
Bud Lite
Michelob
. . .

The answer is a relation with a single attribute, name, and tuples with the name of each beer by Anheuser-Busch, such as Bud.

Meaning of Single-Relation Query

- Begin with the relation in the FROM clause.
- Apply the selection indicated by the WHERE clause.
- Apply the extended projection indicated by the SELECT clause.

Operational Semantics

name	manf
Bud	Anheuser-Busch

Tuple-variable t
loops over all
tuples

Include $t.name$
in the result, if so

Check if
Anheuser-Busch

Operational Semantics --- General

- Think of a *tuple variable* visiting each tuple of the relation mentioned in **FROM**.
- Check if the “current” tuple satisfies the **WHERE** clause.
- If so, compute the attributes or expressions of the **SELECT** clause using the components of this tuple.

* In SELECT clauses

- When there is one relation in the FROM clause, * in the SELECT clause stands for “all attributes of this relation.”
- Example: Using Beers(name, manf):

```
SELECT  *  
FROM    Beers  
WHERE   manf='Anheuser-Busch';
```

Result of Query:

name	manf
Bud	Anheuser-Busch
Bud Lite	Anheuser-Busch
Michelob	Anheuser-Busch
.

Now, the result has each of the attributes of Beers.

Renaming Attributes

- If you want the result to have different attribute names, use "AS <new name>" to rename an attribute.

- Example: Using Beers(name, manf):

```
SELECT name AS beer, manf
FROM Beers
WHERE manf='Anheuser-Busch';
```

Result of Query:

beer	manf
Bud	Anheuser-Busch
Bud Lite	Anheuser-Busch
Michelob	Anheuser-Busch
.

Expressions in SELECT Clauses

- Any expression that makes sense can appear as an element of a SELECT clause.
- **Example:** Using `Sells(bar, beer, price)`:

```
SELECT bar, beer,  
       price*114 AS priceInYen  
FROM Sells;
```

Result of Query

bar	beer	priceInYen
Joe's	Bud	285
Sue's	Miller	342
...

Example: Constants as Expressions

□ Using Likes2(drinker, beer):

```
SELECT drinker,  
       'likes Bud' AS whoLikesBud  
FROM Likes2  
WHERE beer = 'Bud';
```


Result of Query

drinker	whoLikesBud
Sally	likes Bud
Fred	likes Bud
...	...

Example: Information Integration

- We often build “data warehouses” from the data at many “sources.”
- Suppose each bar has its own relation `Menu(beer, price)` .
- To contribute to `Sells(bar, beer, price)` we need to query each bar and insert the name of the bar.

Information Integration --- (2)

□ For instance, at Joe's Bar we can issue the query:

```
SELECT 'Joes Bar', beer, price  
FROM Menu;
```

Complex Conditions in WHERE Clause

- Boolean operators AND, OR, NOT.
- Comparisons =, <>, <, >, <=, >=.
- And many other operators that produce boolean-valued results.

Example: Complex Condition

- Using `Sells(bar, beer, price)`, find the price Joe's Bar charges for Bud:

```
SELECT price
FROM Sells
WHERE bar = 'Joes Bar' AND
       beer = 'Bud';
```

Patterns

- A condition can compare a string to a pattern by:
 - `<Attribute> LIKE <pattern>` or
`<Attribute> NOT LIKE <pattern>`
- *Pattern* is a quoted string with
 - `%` = "any string";
 - `_` = "any character."

Example: LIKE

- Using Drinkers(name, addr, phone) find the drinkers with exchange 555:

```
SELECT name
FROM Drinkers
WHERE phone LIKE '%555-____';
```

NULL Values

- Tuples in SQL relations can have NULL as a value for one or more components.
- Meaning depends on context. Two common cases:
 - *Missing value* : e.g., we know Joe's Bar has some address, but we don't know what it is.
 - *Inapplicable* : e.g., the value of attribute *spouse* for an unmarried person.

Comparing NULL's to Values

- The logic of conditions in SQL is really **3-valued logic**: TRUE, FALSE, **UNKNOWN**.
- Comparing any value (including NULL itself) with NULL yields UNKNOWN.
- A tuple is in a query answer iff the WHERE clause is TRUE (not FALSE or UNKNOWN).

Three-Valued Logic

- To understand how AND, OR, and NOT work in 3-valued logic, think of **TRUE** = 1, **FALSE** = 0, and **UNKNOWN** = $\frac{1}{2}$.
- AND = MIN; OR = MAX, NOT(x) = $1-x$.
- Example:
TRUE AND (FALSE OR NOT(UNKNOWN)) =
MIN(1, MAX(0, (1 - $\frac{1}{2}$))) =
MIN(1, MAX(0, $\frac{1}{2}$)) = MIN(1, $\frac{1}{2}$) = $\frac{1}{2}$.

Surprising Example

□ From the following Sells relation:

bar	beer	price
Joe's Bar	Bud	NULL

SELECT bar

FROM Sells

WHERE price < 2.00 OR price >= 2.00;

← UNKNOWN →

← UNKNOWN →

← UNKNOWN →

Reason: 2-Valued Laws \neq 3-Valued Laws

- Some common laws, like commutativity of AND, hold in 3-valued logic.
- But not others, e.g., the *law of the excluded middle* : $p \text{ OR } \text{NOT } p = \text{TRUE}$.
 - When $p = \text{UNKNOWN}$, the left side is $\text{MAX}(\frac{1}{2}, (1 - \frac{1}{2})) = \frac{1}{2} \neq 1$.

Multirelation Queries

- Interesting queries often combine data from more than one relation.
- We can address several relations in one query by listing them all in the FROM clause.
- Distinguish attributes of the same name by "<relation>.<attribute>".

Example: Joining Two Relations

- Using relations `Likes2(drinker, beer)` and `Frequents(drinker, bar)`, find the beers liked by at least one person who frequents Joe's Bar.

```
SELECT beer
FROM Likes2, Frequents
WHERE bar = 'Joes Bar' AND
      Frequents.drinker =
        Likes2.drinker;
```

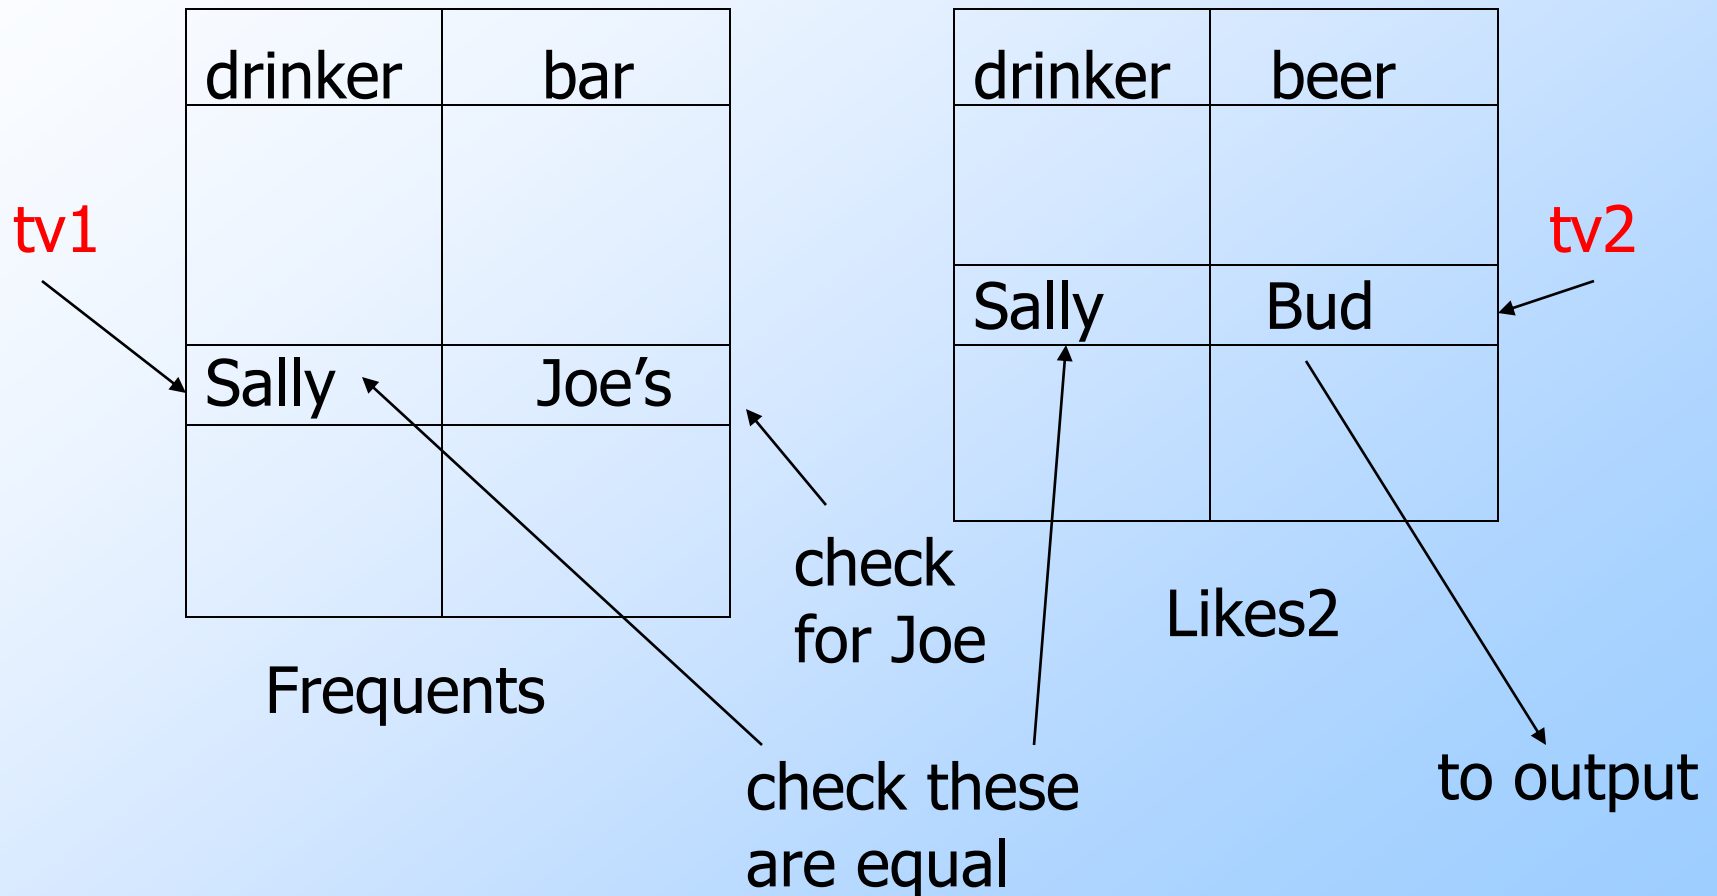
Formal Semantics

- Almost the same as for single-relation queries:
 1. Start with the **product** of all the relations in the **FROM** clause.
 2. Apply the **selection** condition from the **WHERE** clause.
 3. **Project** onto the list of attributes and expressions in the **SELECT** clause.

Operational Semantics

- Imagine **one tuple-variable for each relation** in the FROM clause.
 - These tuple-variables visit each combination of tuples, one from each relation.
- If the tuple-variables are pointing to tuples that satisfy the WHERE clause, send these tuples to the SELECT clause.

Example



Explicit Tuple-Variables

- Sometimes, a query needs to use two copies of the same relation.
- Distinguish copies by following the relation name by the name of a **tuple-variable**, in the **FROM** clause.
- It's always an option to rename relations this way, even when not essential.

Example: Self-Join

- From **Beers(name, manf)**, find all pairs of beers by the same manufacturer.
 - Do not produce pairs like (Bud, Bud).
 - Produce pairs in alphabetic order, e.g. (Bud, Miller), not (Miller, Bud).

```
SELECT b1.name, b2.name
FROM Beers b1, Beers b2
WHERE b1.manf = b2.manf AND
      b1.name < b2.name;
```

Subqueries

- A parenthesized SELECT-FROM-WHERE statement (*subquery*) can be used as a value in a number of places, including **FROM** and **WHERE clauses**.
- **Example**: in place of a relation in the FROM clause, we can use a subquery and then query its result.
 - Must use a tuple-variable to name tuples of the result.

Example: Subquery in FROM


- Find the beers liked by at least one person who frequents Joe's Bar.

```
SELECT beer
```

```
FROM Likes2, (SELECT drinker  
FROM Frequents  
WHERE bar = 'Joes Bar') JD
```

```
WHERE Likes2.drinker=JD.drinker;
```

Drinkers who
frequent Joe's Bar



Subqueries That Return One Tuple

- If a subquery is guaranteed to produce one tuple, then the subquery can be used as a value.
 - Usually, the tuple has one component.
 - A run-time error occurs if there is no tuple or more than one tuple.

Example: Single-Tuple Subquery

- Using `Sells(bar, beer, price)`, find the bars that serve Miller for the same price Joe charges for Bud.
- Two queries would surely work:
 1. Find the price Joe charges for Bud.
 2. Find the bars that serve Miller at that price.

Query + Subquery Solution

SELECT bar

FROM Sells

WHERE beer = 'Miller' AND


price = (SELECT price

FROM Sells

WHERE bar = 'Joes Bar'

AND beer = 'Bud');

The price at
which Joe
sells Bud



The IN Operator

- `<tuple> IN (<subquery>)` is true if and only if the tuple is a member of the relation produced by the subquery.
 - Opposite: `<tuple> NOT IN (<subquery>)`.
- IN-expressions can appear in WHERE clauses.

Example: IN

- Using `Beers(name, manf)` and `Likes2(drinker, beer)`, find the name and manufacturer of each beer that Fred likes.

`SELECT *`

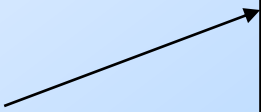
`FROM Beers`

`WHERE name IN (SELECT beer`

`FROM Likes2`

`WHERE drinker = 'Fred');`

The set of
beers Fred
likes



What is the difference?

```
SELECT a  
FROM R, S  
WHERE R.b = S.b;
```

```
SELECT a  
FROM R  
WHERE b IN (SELECT b FROM S);
```

IN is a Predicate About R's Tuples

```
SELECT a
FROM R
WHERE b IN (SELECT b FROM S);
```

Two 2's

One loop, over
the tuples of R

a	b
1	2
3	4

R

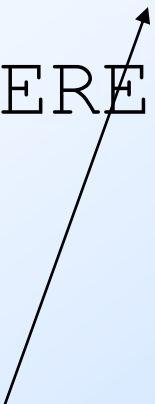
b	c
2	5
2	6

S

(1,2) satisfies
the condition;
1 is output once.

This Query Pairs Tuples from R, S

```
SELECT a
FROM R, S
WHERE R.b = S.b;
```



Double loop, over
the tuples of R and S

a	b
1	2
3	4

R

b	c
2	5
2	6

S

(1,2) with (2,5)
and (1,2) with
(2,6) both satisfy
the condition;
1 is output twice.

The Exists Operator

- **EXISTS**(<subquery>) is true if and only if the subquery result is not empty.
- **Example**: From **Beers(name, manf)** , find those beers that are the unique beer by their manufacturer.

Example: EXISTS

```
SELECT name  
FROM Beers b1  
WHERE NOT EXISTS (
```

Notice scope rule: manf refers to closest nested FROM with a relation having that attribute.

Set of beers with the same manf as b1, but not the same beer

```
SELECT *  
FROM Beers  
WHERE manf = b1.manf AND  
      name <> b1.name);
```

Notice the SQL "not equals" operator

The Operator ANY

- $x = \text{ANY}(\langle \text{subquery} \rangle)$ is a boolean condition that is true iff x equals at least one tuple in the subquery result.
 - $=$ could be any comparison operator.
- **Example:** $x \geq \text{ANY}(\langle \text{subquery} \rangle)$ means x is not the uniquely smallest tuple produced by the subquery.
 - Note tuples must have one component only.

The Operator ALL

- $x <> \text{ALL}(<\text{subquery}>)$ is true iff for every tuple t in the relation, x is not equal to t .
 - That is, x is not in the subquery result.
- $<>$ can be any comparison operator.
- **Example:** $x \geq \text{ALL}(<\text{subquery}>)$ means there is no tuple larger than x in the subquery result.

Example: ALL

- From **Sells(bar, beer, price)**, find the beer(s) sold for the highest price.

SELECT beer

FROM Sells

WHERE price >= ALL(
SELECT price
FROM Sells);

price from the outer
Sells must not be
less than any price.

Union, Intersection, and Difference

□ Union, intersection, and difference of relations are expressed by the following forms, each involving subqueries:

□ (<subquery>) UNION (<subquery>)

□ (<subquery>) INTERSECT (<subquery>)

□ (<subquery>) MINUS (<subquery>)

/* Some DBMSs use EXCEPT instead of MINUS */

Example: Intersection

- Using `Likes2(drinker, beer)`, `Sells(bar, beer, price)`, and `Frequents(drinker, bar)`, find the drinkers and beers such that:
 1. The drinker likes the beer, and
 2. The drinker frequents at least one bar that sells the beer.

Notice trick:
subquery is
really a stored
table.

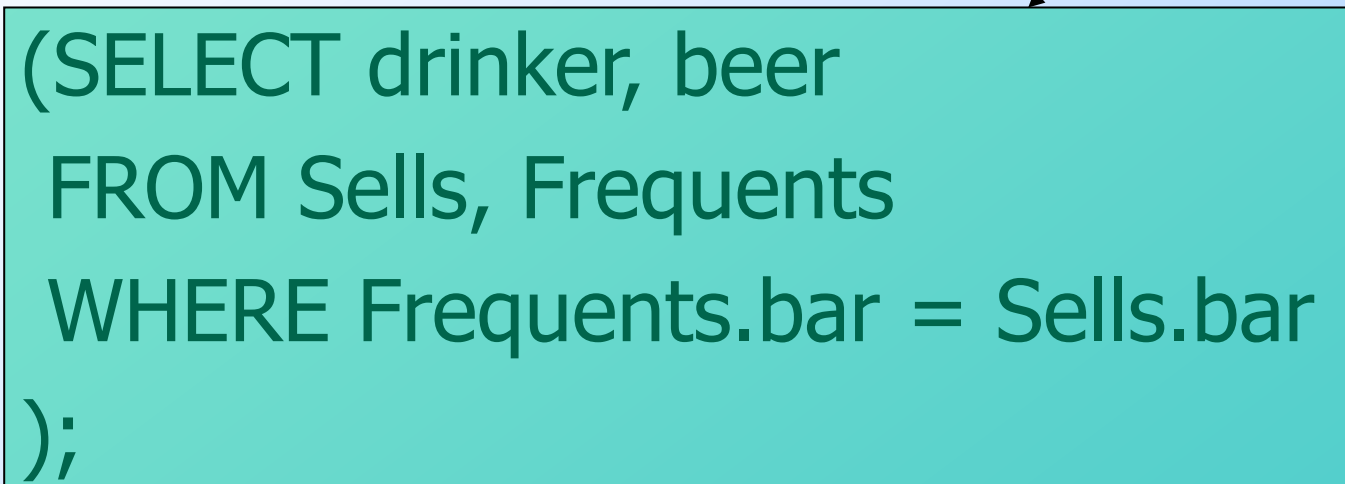
Solution



```
(SELECT * FROM Likes2)
```

INTERSECT

The drinker frequents
a bar that sells the
beer.



```
(SELECT drinker, beer  
FROM Sells, Frequents  
WHERE Frequents.bar = Sells.bar  
);
```

Bag Semantics

- Although the SELECT-FROM-WHERE statement uses bag semantics, the default for union, intersection, and difference is set semantics.
- That is, duplicates are eliminated as the operation is applied.

Motivation: Efficiency

- When doing projection, it is easier to avoid eliminating duplicates.
 - Just work tuple-at-a-time.
- For intersection or difference, it is most efficient to sort the relations first.
 - At that point you may as well eliminate the duplicates anyway.

Controlling Duplicate Elimination

- Force the result to be a set by
SELECT **DISTINCT** . . .
- Force the result to be a bag (i.e., don't eliminate duplicates) by ALL, as in
. . . **UNION ALL** . . .

Example: DISTINCT

- From `Sells(bar, beer, price)`, find all the different prices charged for beers:

```
SELECT DISTINCT price  
FROM Sells;
```

- Notice that without `DISTINCT`, each price would be listed as many times as there were bar/beer pairs at that price.

Example: ALL

- Using relations **Frequents(drinker, bar)** and **Likes2(drinker, beer)**:

```
(SELECT drinker FROM Frequents)
  MINUS ALL /* Oracle doesn't support */
(SELECT drinker FROM Likes2);
```

- Lists drinkers who frequent more bars than they like beers, and does so as many times as the difference of those counts.

Join Expressions

- SQL provides several versions of (bag) joins.
- These expressions can be stand-alone queries or used in place of relations in a FROM clause.

Products and Natural Joins

- Natural join:

R **NATURAL JOIN** S;

- Product:

R **CROSS JOIN** S;

- **Example:**

Likes2 **NATURAL JOIN** Sells;

- Relations can be **parenthesized subqueries**, as well.

Theta Join

- $R \text{ JOIN } S \text{ ON } \langle \text{condition} \rangle$
- **Example:** using $\text{Drinkers}(\text{name}, \text{addr})$ and $\text{Frequents}(\text{drinker}, \text{bar})$:

```
Drinkers JOIN Frequents ON  
    name = drinker;
```

gives us all (d, a, d, b) quadruples such that drinker d lives at address a and frequents bar b .