

Views, Indexes

Virtual and Materialized Views
Speeding Accesses to Data

Views

- A *view* is a relation defined in terms of stored tables (called *base tables*) and other views.
- Two kinds:
 1. *Virtual* = not stored in the database; just a query for constructing the relation.
 2. *Materialized* = actually constructed and stored.

Declaring Views

- Declare by:

```
CREATE [MATERIALIZED] VIEW  
    <name> AS <query>;
```

- Default is virtual.

Example: View Definition

- **CanDrink(drinker, beer)** is a view “containing” the drinker-beer pairs such that the drinker frequents at least one bar that serves the beer:

```
CREATE VIEW CanDrink AS
  SELECT drinker, beer
  FROM Frequents, Sells
  WHERE Frequents.bar = Sells.bar;
```

Example: Accessing a View

- Query a view **as if it were a base table**.
 - **Also**: a limited ability to **modify** views if it makes sense as a modification of one underlying base table.
- **Example query**:

```
SELECT beer FROM CanDrink  
WHERE drinker = 'Sally';
```

Triggers on Views

- Generally, it is impossible to modify a virtual view, because it doesn't exist.
- But an **INSTEAD OF trigger** lets us interpret view modifications in a way that makes sense.
- **Example:** View Synergy has (**drinker, beer, bar**) triples such that the bar serves the beer, the drinker frequents the bar and likes the beer.

Example: The View

CREATE VIEW Synergy AS

SELECT Likes.drinker, Likes.beer, Sells.bar

Pick one copy of
each attribute

FROM Likes, Sells, Frequents

WHERE Likes.drinker = Frequents.drinker

AND Likes.beer = Sells.beer

AND Sells.bar = Frequents.bar;

Natural join of Likes,
Sells, and Frequents

Interpreting a View Insertion

- We cannot insert into Synergy --- it is a virtual view.
- But we can use an **INSTEAD OF** trigger to turn a **(drinker, beer, bar)** triple into **three insertions** of projected pairs, one for each of Likes, Sells, and Frequents.
 - Sells.price will have to be NULL.

The Trigger

```
CREATE TRIGGER ViewTrig
  INSTEAD OF INSERT ON Synergy
  REFERENCING NEW ROW AS n
  FOR EACH ROW
  BEGIN
    INSERT INTO LIKES VALUES(n.drinker, n.beer);
    INSERT INTO SELLS(bar, beer) VALUES(n.bar, n.beer);
    INSERT INTO FREQUENTS VALUES(n.drinker, n.bar);
  END;
```

Materialized Views

- **Problem:** each time a base table changes, the materialized view may change.
 - Cannot afford to recompute the view with each change.
- **Solution:** **Periodic reconstruction** of the materialized view, which is otherwise “out of date.”

Example: Axxess/Class Mailing List

- The class mailing list `cs145-aut0708-students` is in effect a materialized view of the class enrollment in Axxess.
- Actually updated four times/day.
 - You can enroll and miss an email sent out after you enroll.

Example: A Data Warehouse

- Wal-Mart stores every sale at every store in a database.
- **Overnight**, the sales for the day are used to **update** a *data warehouse* = materialized views of the sales.
- The warehouse is used by analysts to predict trends and move goods to where they are selling best.

Indexes

- *Index* = data structure used to speed access to tuples of a relation, given values of one or more attributes.
- Could be a **hash table**, but in a DBMS it is always a balanced search tree with giant nodes (a full disk page) called a *B-tree*.

Declaring Indexes

- No standard!

- Typical syntax:

```
CREATE INDEX BeerInd ON  
  Beers (manf) ;
```

```
CREATE INDEX SellInd ON  
  Sells (bar, beer) ;
```

Using Indexes

- Given a value v , the index takes us to only those tuples that have v in the attribute(s) of the index.
- **Example:** use BeerInd and SellInd to find the prices of beers manufactured by Pete's and sold by Joe. (next slide)

Using Indexes --- (2)

```
SELECT price FROM Beers, Sells  
WHERE manf = 'Petes' AND  
       Beers.name = Sells.beer AND  
       bar = 'Joes Bar';
```

1. Use **BeerInd** to get all the beers made by Pete's.
2. Then use **SellInd** to get prices of those beers, with bar = 'Joes Bar'

Database Tuning

- A major problem in making a database run fast is deciding which indexes to create.
- Pro: An index speeds up queries that can use it.
- Con: An index slows down all modifications on its relation because the index must be modified too.

Example: Tuning

- Suppose the only things we did with our beers database was:
 1. Insert new facts into a relation (10%).
 2. Find the price of a given beer at a given bar (90%).
- Then **SellInd** on Sells(bar, beer) would be wonderful, but **BeerInd** on Beers(manf) would be harmful.

Tuning Advisors

- A major research thrust.
 - Because hand tuning is so hard.
- An advisor gets a *query load*, e.g.:
 1. Choose random queries from the history of queries run on the database, or
 2. Designer provides a sample workload.

Tuning Advisors --- (2)

- The advisor generates candidate indexes and evaluates each on the workload.
 - Feed each sample query to the query optimizer, which assumes only this one index is available.
 - Measure the improvement/degradation in the average running time of the queries.